

# Step 1: Our Stakeholder and the Business Problem.

## HGTV SHOW

-HGTV is launching a show based on homes in kings county. The goal of the show is to show homeowners what predictors lead to a higher listing price so the owner can sell and maximize profit. The goal is getting homes with good grades and base quality that may need a couple of alterations to maximize it's value.

## IMPORTING LIBRARIES

```
In [1]: #importing the necessary libraries for the project.
import pandas as pd
import numpy as np
import seaborn as sns
from matplotlib import pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.feature_selection import RFE
from sklearn.preprocessing import PolynomialFeatures, StandardScaler
import statsmodels
from statsmodels.formula.api import ols
from sklearn.model_selection import train_test_split
from sklearn.dummy import DummyRegressor
from statsmodels.tools.eval_measures import rmse
from statsmodels.api import qqplot
from scipy import stats
from sklearn.preprocessing import OneHotEncoder
from folium.plugins import FastMarkerCluster
import folium
from sklearn.metrics import r2_score
import branca.colormap as cm
```

```
In [2]: #loading in the data and making a pandas dataframe.
df = pd.read_csv('data/kc_house_data.csv')
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     21597 non-null  int64
1   date                   21597 non-null  object
2   price                  21597 non-null  float64
3   bedrooms               21597 non-null  int64
4   bathrooms              21597 non-null  float64
5   sqft_living            21597 non-null  int64
6   sqft_lot               21597 non-null  int64
7   floors                 21597 non-null  float64
8   waterfront             19221 non-null  object
9   view                   21534 non-null  object
10  condition              21597 non-null  object
11  grade                  21597 non-null  object
12  sqft_above             21597 non-null  int64
13  sqft_basement          21597 non-null  object
14  yr_built               21597 non-null  int64
15  yr_renovated           17755 non-null  float64
16  zipcode                21597 non-null  int64
17  lat                    21597 non-null  float64
18  long                   21597 non-null  float64
19  sqft_living15          21597 non-null  int64
20  sqft_lot15             21597 non-null  int64
dtypes: float64(6), int64(9), object(6)
memory usage: 3.5+ MB

```

## EDA(EXPLORITORY DATA ANALYSIS)

```

In [3]: #Exploring the columns and values.
df.head()

```

```

Out[3]:

```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
0	7129300520	10/13/2014	221900.0	3	1.00	1180	5650	1.0	
1	6414100192	12/9/2014	538000.0	3	2.25	2570	7242	2.0	
2	5631500400	2/25/2015	180000.0	2	1.00	770	10000	1.0	
3	2487200875	12/9/2014	604000.0	4	3.00	1960	5000	1.0	
4	1954400510	2/18/2015	510000.0	3	2.00	1680	8080	1.0	

5 rows × 21 columns

-Exploring some low hanging data. some standouts: Average house in this county have 3 bedrooms and 2 bathrooms. -Theres a hosue in the set that gas 33 bedrooms and 8 bathroom which will defiently be an outlier. -Oldest house is dated to 1900. While he most recent hosue was built in 2015.

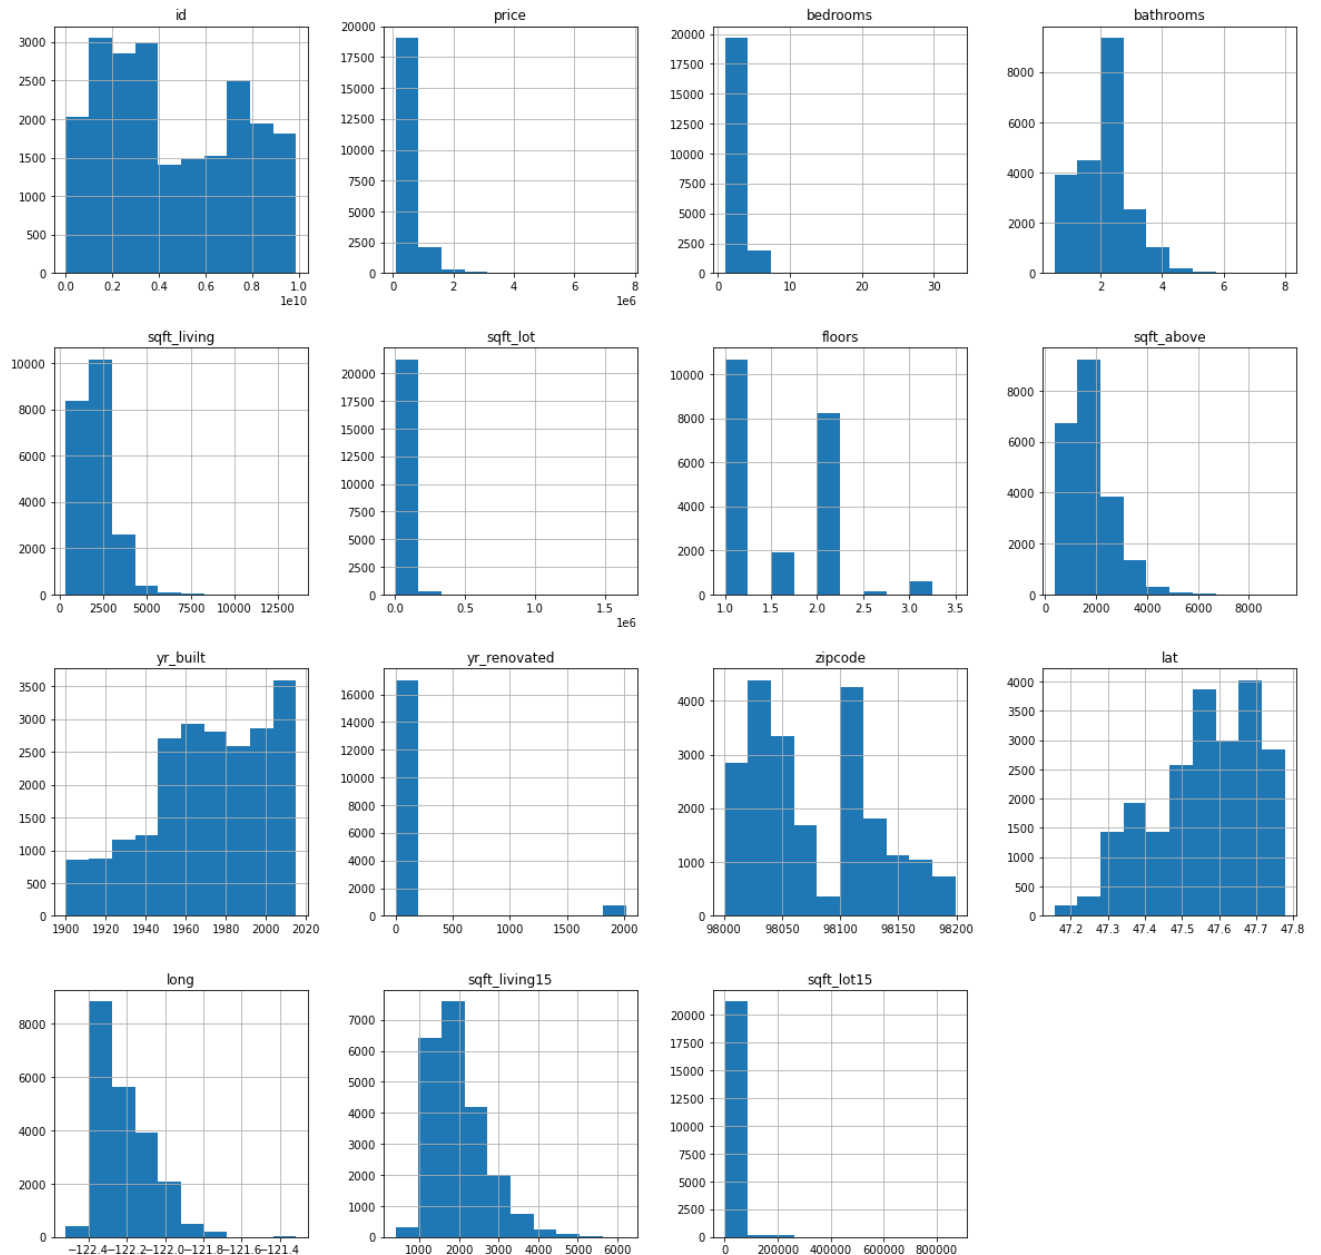
In [6]: `df.describe()`

Out[6]:

		id	price	bedrooms	bathrooms	sqft_living	sqft_lot
<b>count</b>	2.159700e+04	2.159700e+04	21597.000000	21597.000000	21597.000000	2.159700e+04	
<b>mean</b>	4.580474e+09	5.402966e+05	3.373200	2.115826	2080.321850	1.509941e+04	
<b>std</b>	2.876736e+09	3.673681e+05	0.926299	0.768984	918.106125	4.141264e+04	
<b>min</b>	1.000102e+06	7.800000e+04	1.000000	0.500000	370.000000	5.200000e+02	
<b>25%</b>	2.123049e+09	3.220000e+05	3.000000	1.750000	1430.000000	5.040000e+03	
<b>50%</b>	3.904930e+09	4.500000e+05	3.000000	2.250000	1910.000000	7.618000e+03	
<b>75%</b>	7.308900e+09	6.450000e+05	4.000000	2.500000	2550.000000	1.068500e+04	
<b>max</b>	9.900000e+09	7.700000e+06	33.000000	8.000000	13540.000000	1.651359e+06	

Running a histogram of the dataframe. So visibily see the distribution. Apart from Latitude and Year Built. Most of the variables distrubte the right.

In [5]: `df.hist(figsize=(20,20));`



-Dropping houses that are repeated. some houses are repeated if they had renovations done. -Filling in waterfront's NA values with 0 -Dropping years renovated column because a large amount of 0s

```
In [7]: #exploring and filtering

df.drop_duplicates(subset='id',keep='first',inplace=True)
df['waterfront'].fillna(0, inplace=True)
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 21420 entries, 0 to 21596
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    21420 non-null  int64
1   date                 21420 non-null  object
2   price               21420 non-null  float64
3   bedrooms            21420 non-null  int64
4   bathrooms           21420 non-null  float64
5   sqft_living         21420 non-null  int64
6   sqft_lot            21420 non-null  int64
7   floors              21420 non-null  float64
8   waterfront          21420 non-null  object
9   view               21357 non-null  object
10  condition            21420 non-null  object
11  grade               21420 non-null  object
12  sqft_above          21420 non-null  int64
13  sqft_basement       21420 non-null  object
14  yr_built            21420 non-null  int64
15  yr_renovated        17616 non-null  float64
16  zipcode             21420 non-null  int64
17  lat                 21420 non-null  float64
18  long                21420 non-null  float64
19  sqft_living15       21420 non-null  int64
20  sqft_lot15          21420 non-null  int64
dtypes: float64(6), int64(9), object(6)
memory usage: 3.6+ MB

```

```

In [8]: for column in list(df.columns):
        print(column, sum(df[column].isnull()))

```

```

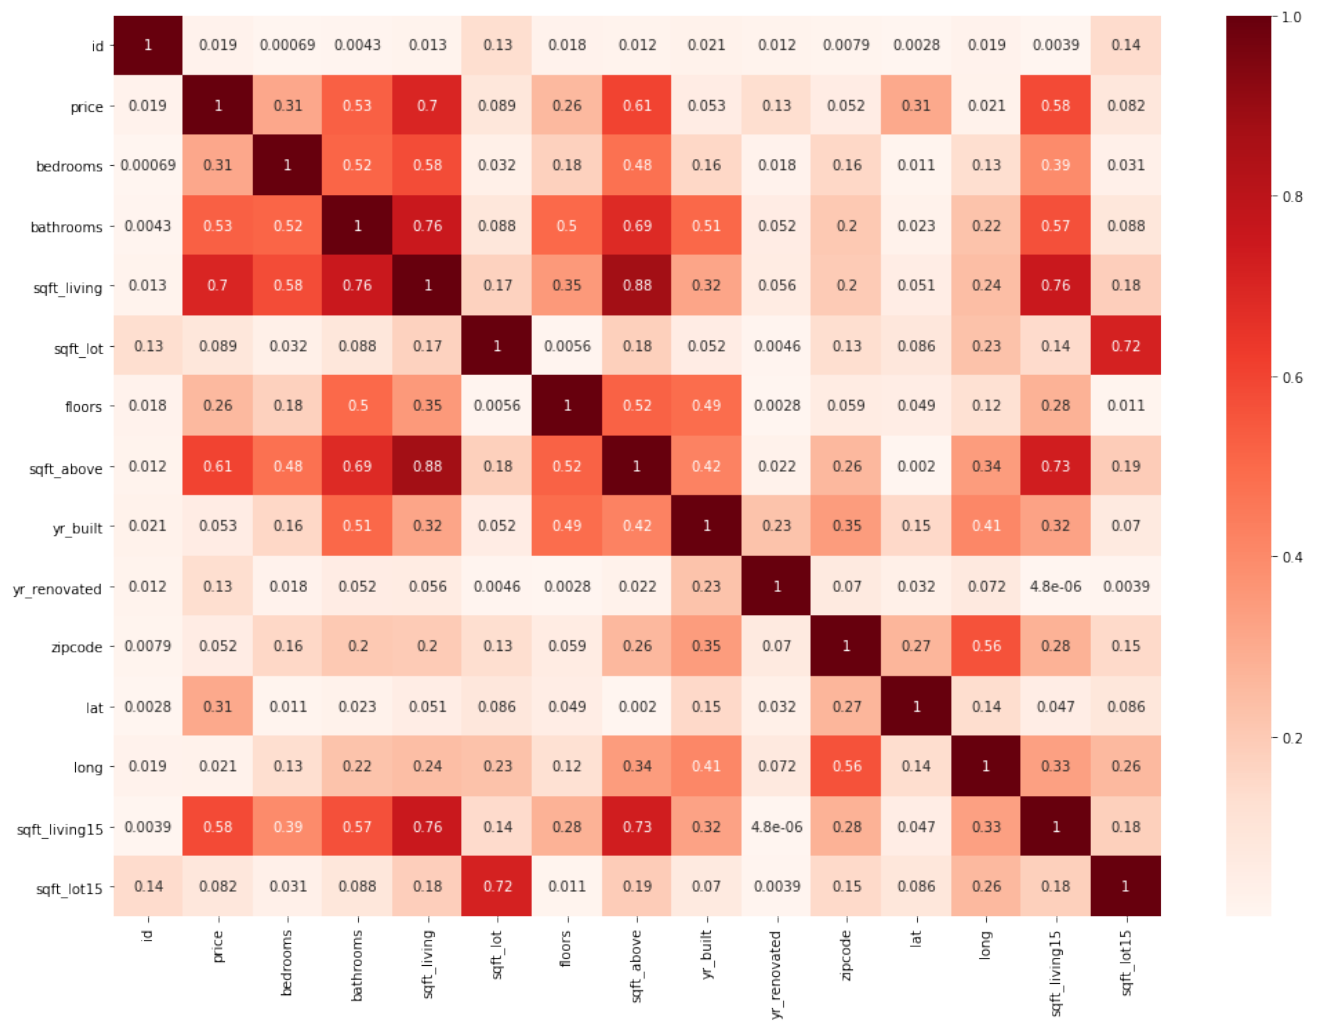
id 0
date 0
price 0
bedrooms 0
bathrooms 0
sqft_living 0
sqft_lot 0
floors 0
waterfront 0
view 63
condition 0
grade 0
sqft_above 0
sqft_basement 0
yr_built 0
yr_renovated 3804
zipcode 0
lat 0
long 0
sqft_living15 0
sqft_lot15 0

```

# Pre Model Assumptions

```
In [12]: corr = df.corr().abs()
fig, ax=plt.subplots(figsize=(17,12))
fig.suptitle('Correlations', fontsize=30, y=.95, fontname='Silom')
heatmap = sns.heatmap(corr, cmap='Reds', annot=True)
```

## Correlations



```
In [13]: features = []
correlations = []
for idx, correlation in corr['price'].T.iteritems():
    if correlation >= .30 and idx != 'price':
        features.append(idx)
        correlations.append(correlation)
corr_with_price = pd.DataFrame({'Correlations':correlations, 'Features': feat

Multicollinear_Features = []
Multicollinear_Corr = []
def check_multicollinearity(feature):
    for idx, correlation in corr[feature].T.iteritems():
        if correlation >= .80 and idx != feature:
            Multicollinear_Features.append([feature, idx])
            Multicollinear_Corr.append(correlation)

for feature in corr:
    check_multicollinearity(feature)
MC_df = pd.DataFrame({'Correlations':Multicollinear_Corr, 'Features': Multico
print('Multicollinear Features')
display(MC_df)
print('Correlations with Price')
display(MC_df)
```

Multicollinear Features

	Correlations	Features
0	0.876533	[sqft_living, sqft_above]
1	0.876533	[sqft_above, sqft_living]

Correlations with Price

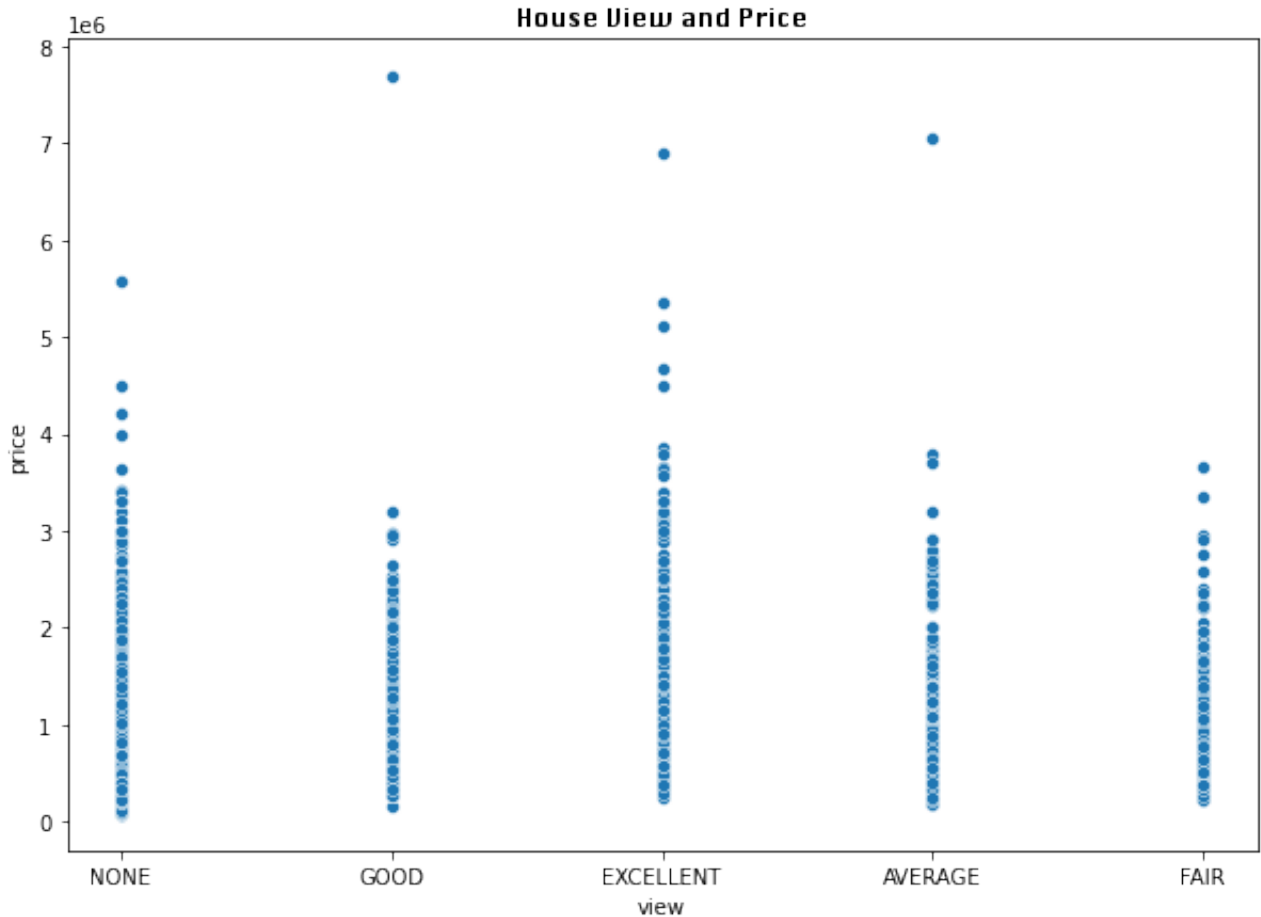
	Correlations	Features
0	0.876533	[sqft_living, sqft_above]
1	0.876533	[sqft_above, sqft_living]

```
In [14]: plt.figure(figsize=(10,7))
sns.scatterplot(df['view'], df['price'])
plt.title('House View and Price', fontsize=12, fontname='silom')
```

/Users/olamideholayinka/opt/anaconda3/envs/learn-env/lib/python3.8/site-packages/seaborn/\_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(

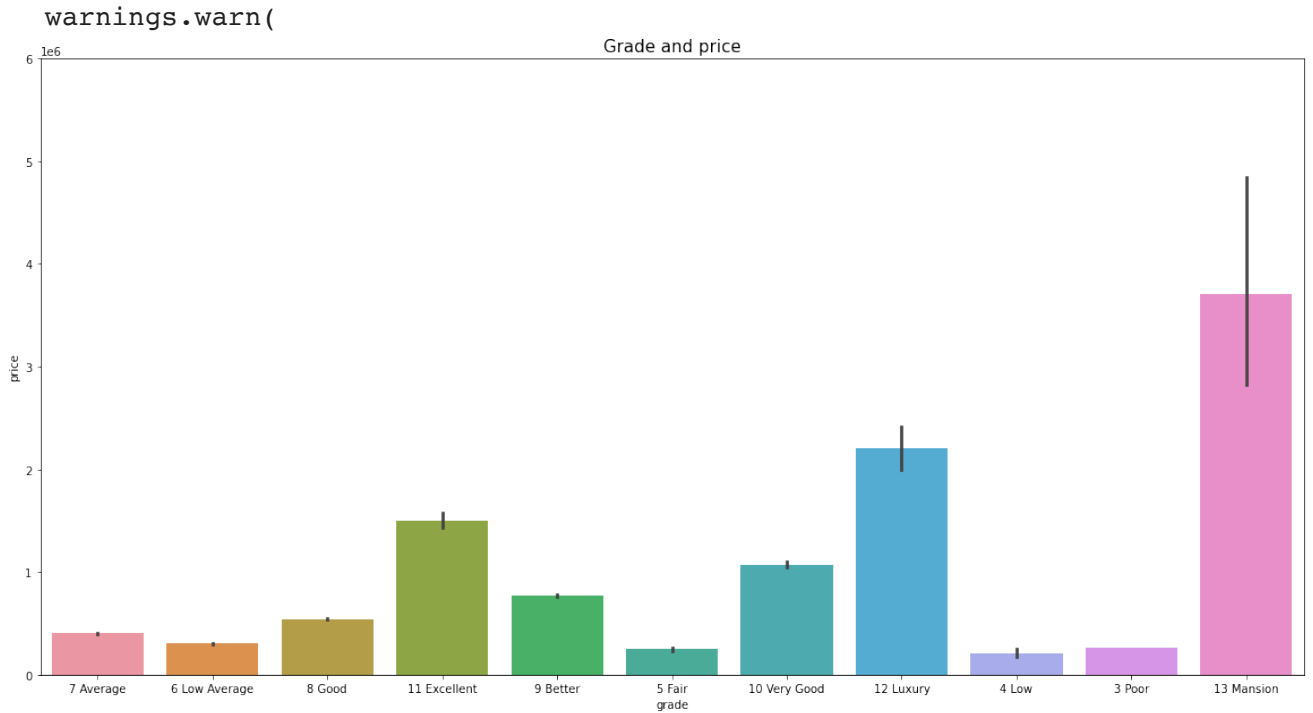
Out[14]: Text(0.5, 1.0, 'House View and Price')



```
In [13]: fig, ax = plt.subplots(figsize=(20,10))
y = df['price']
x = df['grade']
sns.barplot(x,y)
ax.set_ylim(bottom=0,top=6000000)
plt.title('Grade and price', fontsize=15, fontname='silom');
```

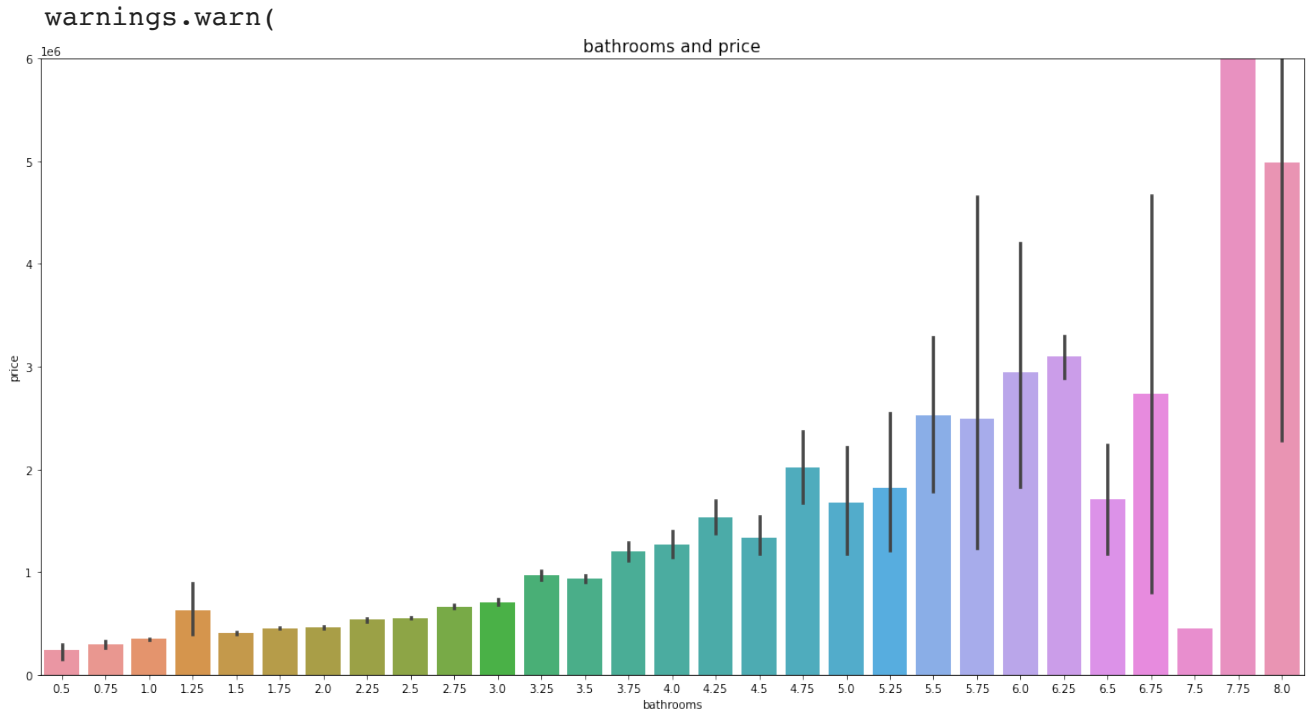


C:\Users\ilene\anaconda3\envs\learn-env\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.



```
In [30]: fig, ax = plt.subplots(figsize =(20,10))
y = df['price']
x = df['bathrooms']
sns.barplot(x,y)
ax.set_ylim(bottom=0,top=6000000)
plt.title('bathrooms and price', fontsize=15, fontname='silom');
```

C:\Users\ilene\anaconda3\envs\learn-env\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.



```
In [15]: def get_center_latlong(df):
           centerlat = (df['lat'].max() + df['lat'].min()) / 2
           centerlong = (df['long'].max() + df['long'].min()) / 2
           return centerlat, centerlong
```

```
In [16]: center = get_center_latlong(df)
```

```
In [17]: colormap = cm.LinearColormap(colors=['darkblue', 'orange', 'yellow', 'lightyellow'])

sqft_map = folium.Map(location=center, zoom_start=10)
for i in range(len(df)):
    folium.Circle(
        location=[df.iloc[i]['lat'], df.iloc[i]['long']],
        radius=10,
        fill=True,
        color=colormap(df.iloc[i]['sqft_living']),
        fill_opacity=0.2).add_to(sqft_map)

sqft_map.add_child(colormap)
```

Out[17]: Make this Notebook Trusted to load map: File -> Trust Notebook

```
In [35]: df = df[(df["bathrooms"] == 7.5)]
df
```

```
Out[35]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	wa
<b>8537</b>	424049043	8/11/2014	450000.0	9	7.5	4050	6504	2.0	

1 rows × 21 columns

```
In [37]: df.sqft_living.describe()
```



```
Out[37]: count      1.0
mean      4050.0
std         NaN
min      4050.0
25%      4050.0
50%      4050.0
75%      4050.0
max      4050.0
Name: sqft_living, dtype: float64
```


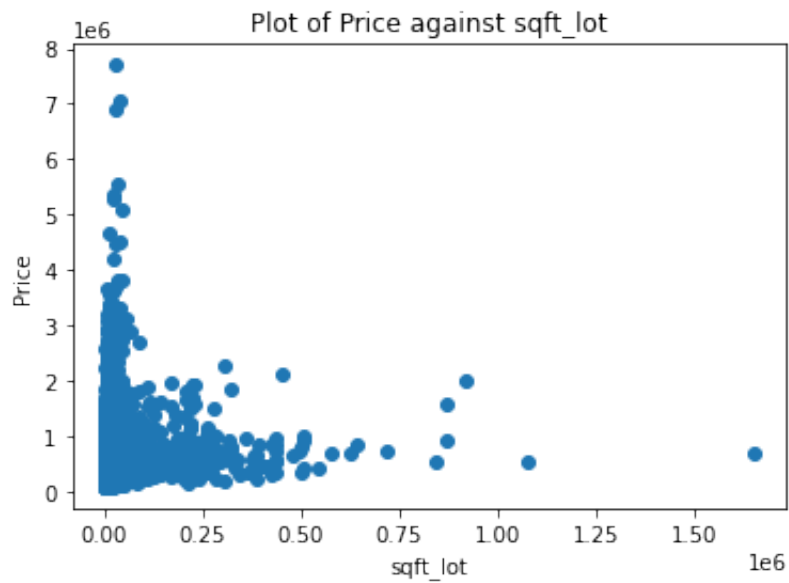
```
In [ ]: # CHECKING FOR LINEARITY
```



```
In [22]: features = X_train.columns

for x in features:
    plt.scatter(X_train[x], y_train)
    plt.title(f'Plot of Price against {x}')
    plt.xlabel(x)
    plt.ylabel('Price')
    plt.show()

# also plot sales against itself
plt.scatter(y_train.index, y_train)
plt.hlines(y_train.mean(), 0, 200)
plt.xlabel('Index Value')
plt.ylabel('Price')
plt.title('Variance of Price')
plt.show()
```

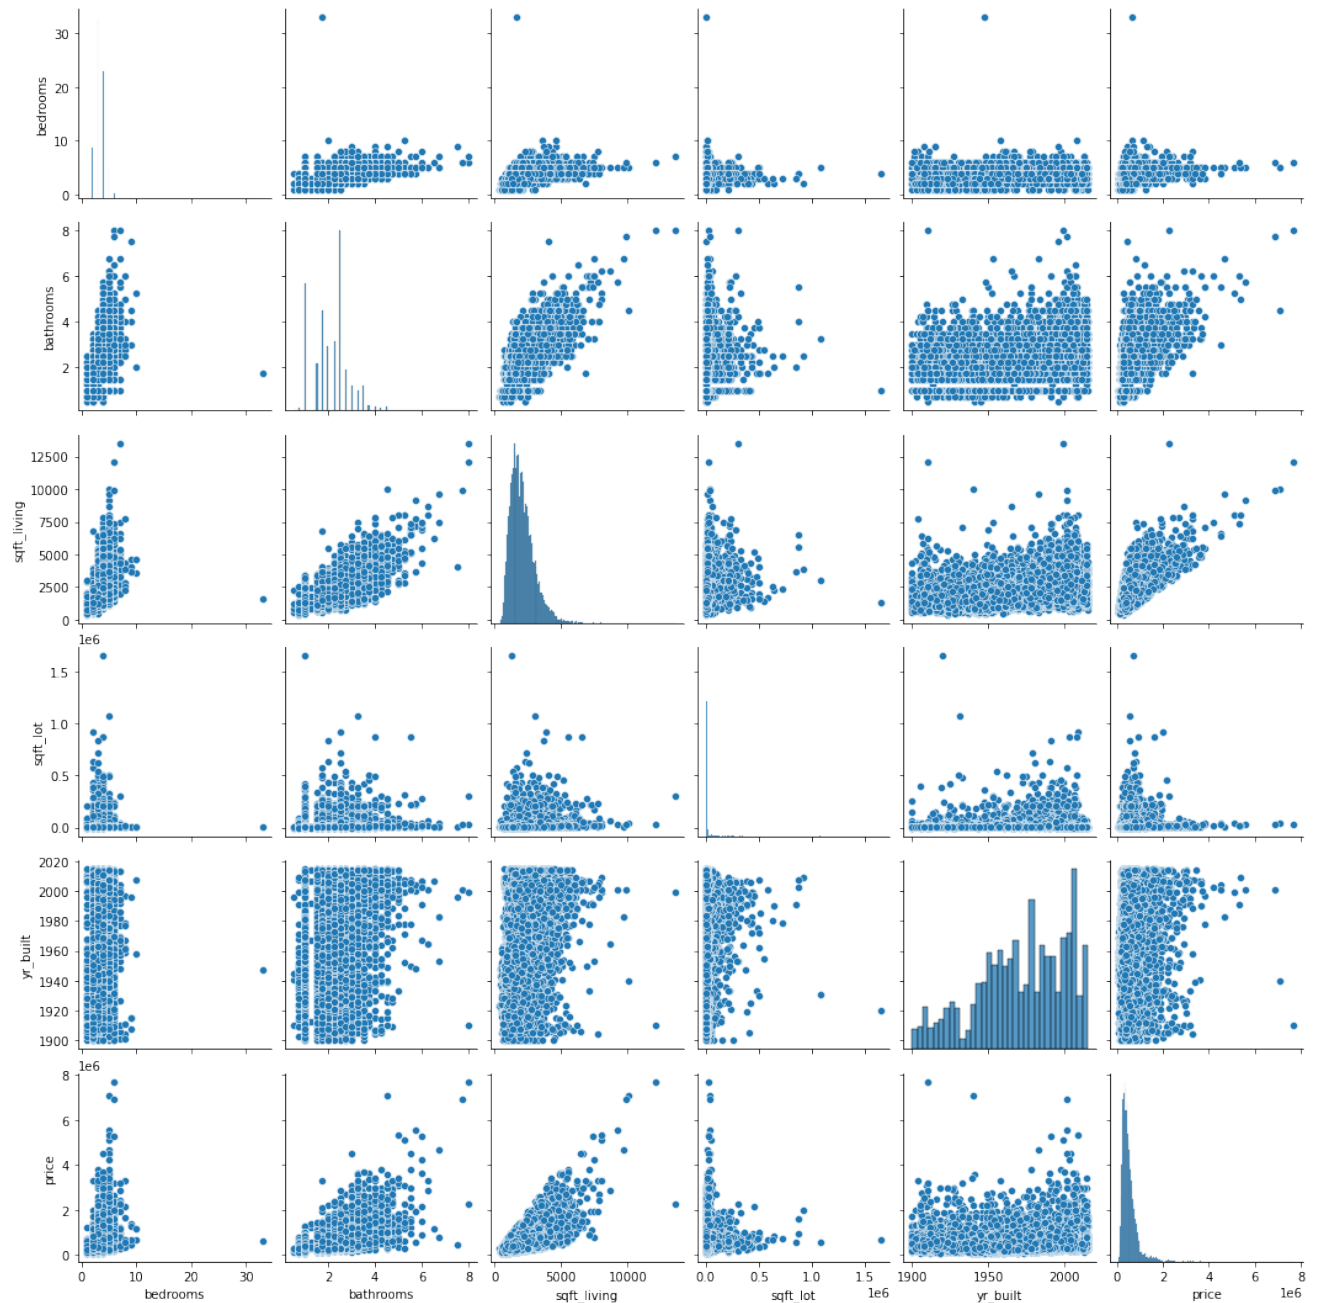






```
In [23]: train_df = pd.concat([X_train, y_train], axis =1)
```

```
In [24]: sns.pairplot(train_df)
plt.show()
```



## STEP 3 TRAIN TEST SPLIT

```
In [21]: y = df.price
X = df[['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'yr_built', 'grade']
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.20, random_state=42)
```

```
In [ ]:
```

## STEP 4 ONE HOT ENCODE CATEGORICAL VARIABLE

```
In [ ]: #grade
```

```
In [25]: ohe = OneHotEncoder()  
ohe.fit_transform(X_train)
```

```
Out[25]: <17136x9494 sparse matrix of type '<class 'numpy.float64'>'  
         with 102816 stored elements in Compressed Sparse Row format>
```

```
In [ ]:
```

```
In [26]: train_df.corr()
```

```
Out[26]:
```

	bedrooms	bathrooms	sqft_living	sqft_lot	yr_built	price
bedrooms	1.000000	0.512126	0.575100	0.040754	0.152713	0.308335
bathrooms	0.512126	1.000000	0.756116	0.091503	0.503744	0.526030
sqft_living	0.575100	0.756116	1.000000	0.180099	0.314592	0.702331
sqft_lot	0.040754	0.091503	0.180099	1.000000	0.050806	0.092713
yr_built	0.152713	0.503744	0.314592	0.050806	1.000000	0.052383
price	0.308335	0.526030	0.702331	0.092713	0.052383	1.000000

## STEP 5 CREATING DUMMY BASELINE

```
In [27]: #Dummy Regressor for the baseline. DUMMY IS ONLY APPLIED ON TRAINING SET  
dummy_regr = DummyRegressor(strategy='mean')  
dummy_regr.fit(X_train, y_train)  
dummy_regr.predict(X_train)
```

```
Out[27]: array([542950.3173436, 542950.3173436, 542950.3173436, ...,  
                542950.3173436, 542950.3173436, 542950.3173436])
```

```
In [28]: dummy_regr.score(X_train, y_train)# the R score is zero. The mean of Y expla
```

```
Out[28]: 0.0
```

```
In [ ]: y_train.shape  
y_train
```



# RUNNING A SIMPLE MODEL WITH VARIABLES []. THESE VARIABLES WERE CHOSEN BASED ON RFE OR VARS THAT ARE HIGHLY CORRELATED WITH SALES.

```
In [ ]: # NOW CREATE FIRST SIMPLE MODEL THAT DOES BETTER THAN DUMMY REGRESSOR MEAN.
#EXAMPLES; RECURSIVE FEATURE ELIMINATION. LOOK AT HIGH CORRELATIONS VARIABLES

#MODEL 1
```

```
In [30]: y = df.price
X = df[['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'yr_built', 'grade',
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.20, random_state=42)
```

```
In [31]: dummy_regr = DummyRegressor(strategy='mean')
dummy_regr.fit(X_train, y_train)
dummy_regr.predict(X_train)
```

```
Out[31]: array([542950.3173436, 542950.3173436, 542950.3173436, ...,
542950.3173436, 542950.3173436, 542950.3173436])
```

```
In [32]: dummy_regr.score(X_train, y_train)
```

```
Out[32]: 0.0
```

```
In [33]: grade_trial = pd.DataFrame()
grade_test = pd.DataFrame()
```

```
In [34]: ohe = OneHotEncoder(sparse=False, handle_unknown='ignore')
grade_trial['grade'] = X_train['grade']
grade_trial['ID'] = len(X_train.values)
ohe_data = pd.DataFrame(ohe.fit_transform(grade_trial))
#grade_trial
ohe_data.columns = ohe.get_feature_names(['grade', 'ID'])
ohe_data = ohe_data.drop('ID_17136', axis=1)
#ohe_data

grade_test['grade'] = X_test['grade']
grade_test['ID'] = len(X_test.values)
ohe_test = pd.DataFrame(ohe.transform(grade_test))
ohe_test.columns = ohe.get_feature_names(['grade', 'ID'])
ohe_test = ohe_test.drop('ID_17136', axis=1)
```

```
In [ ]:
```

```
In [35]: np.random.seed(42)

y = df.price
x = df[['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'yr_built', 'grade',

# Train test split with random_state=42 and test_size=0.2

X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.2,
                                                    random_state=42)

# Scale appropriately on TRAINING DATA

ss = StandardScaler()
ss.fit(X_train.drop('grade', axis=1))
Testpreds = ss.transform(X_test.drop('grade', axis=1))
X_preds_st_scaled = ss.transform(X_train.drop('grade', axis=1))

lr = LinearRegression()

# fit and score the model (checkout the test set if there is time)
lr.fit(X_preds_st_scaled, y_train)
lr.coef_
lr.score(X_preds_st_scaled, y_train)
```

Out[35]: 0.5547264615853337

```
In [36]: combined_x = ohe_data.join(pd.DataFrame(X_preds_st_scaled))
combined_test = ohe_test.join(pd.DataFrame(Testpreds))
```

```
In [37]: type(combined_test)
```

Out[37]: pandas.core.frame.DataFrame

```
In [38]: lr = LinearRegression()

# fit and score the model (checkout the test set if there is time)
lr.fit(combined_x.values, y_train)
lr.coef_
lr.score(combined_x.values, y_train)
```

Out[38]: 0.6479566533507228

```
In [39]: model1 = lr.predict(combined_test.values)
lr.coef_
lr.score(combined_test.values, y_test)
```

Out[39]: 0.6313406870043061

```
In [40]: def sk_metrics(y, model):
    from sklearn.metrics import mean_squared_error, mean_absolute_error

    print("Metrics:")
    # R2
    print(f"R2: {r2_score(y, model):.3f}")
    # MAE
    print(f"Mean Absolute Error: {mean_absolute_error(y, model):.3f}")
    # MSE
    print(f"Mean Squared Error: {mean_squared_error(y, model):.3f}")
    # RMSE - just MSE but set squared=False
    print(f"Root Mean Squared Error: {mean_squared_error(y, model, squared=False):.3f}")
    return
```

```
In [41]: sk_metrics(y_test, model1)
```

```
Metrics:
R2: 0.631
Mean Absolute Error: 136819.122
Mean Squared Error: 44643315242.166
Root Mean Squared Error: 211289.648
```

```
In [42]: display(X_train.head())
display(X_test.head())
```

	bedrooms	bathrooms	sqft_living	sqft_lot	yr_built	grade	floors
<b>13634</b>	4	2.50	2240	4616	2001	7 Average	2.0
<b>13631</b>	2	2.25	2470	7658	1954	8 Good	1.0
<b>5717</b>	5	3.50	3530	218472	1999	7 Average	2.0
<b>15431</b>	4	1.50	1580	10230	1945	6 Low Average	1.0
<b>17294</b>	4	2.50	2070	10244	1969	8 Good	1.0

	bedrooms	bathrooms	sqft_living	sqft_lot	yr_built	grade	floors
<b>6132</b>	3	1.00	1020	6000	1900	7 Average	1.5
<b>8993</b>	3	1.75	1620	5500	1950	7 Average	1.0
<b>559</b>	4	2.50	2050	10271	1998	7 Average	2.0
<b>11931</b>	3	2.50	2510	5186	2001	7 Average	2.0
<b>15176</b>	2	2.00	1390	12530	1959	8 Good	1.0

```
In [43]: # SECOND MODEL .R squared is lower than our first.
```

```
In [64]: Model_2 = ols(formula="price ~bathrooms + bedrooms", data=df).fit()
Model_2.summary()
```

Out[64]:

## OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.279
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.279
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	4142.
<b>Date:</b>	Thu, 17 Feb 2022	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	15:32:08	<b>Log-Likelihood:</b>	-3.0140e+05
<b>No. Observations:</b>	21420	<b>AIC:</b>	6.028e+05
<b>Df Residuals:</b>	21417	<b>BIC:</b>	6.028e+05
<b>Df Model:</b>	2		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	-3.547e+04	8361.541	-4.242	0.000	-5.19e+04	-1.91e+04
<b>bathrooms</b>	2.39e+05	3240.812	73.754	0.000	2.33e+05	2.45e+05
<b>bedrooms</b>	2.07e+04	2692.095	7.691	0.000	1.54e+04	2.6e+04

<b>Omnibus:</b>	17197.290	<b>Durbin-Watson:</b>	1.969
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	897026.537
<b>Skew:</b>	3.474	<b>Prob(JB):</b>	0.00
<b>Kurtosis:</b>	33.932	<b>Cond. No.</b>	17.0

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [65]: Model_3 = ols(formula="price ~sqft_living + sqft_living15 + sqft_lot ", data=
Model_3.summary()
#better r squared than our second but not better than our first
```

Out[65]:

## OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.501
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.500
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	7155.
<b>Date:</b>	Thu, 17 Feb 2022	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	15:32:08	<b>Log-Likelihood:</b>	-2.9747e+05
<b>No. Observations:</b>	21420	<b>AIC:</b>	5.949e+05
<b>Df Residuals:</b>	21416	<b>BIC:</b>	5.950e+05
<b>Df Model:</b>	3		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	-1.004e+05	5456.168	-18.400	0.000	-1.11e+05	-8.97e+04
<b>sqft_living</b>	245.0885	2.970	82.528	0.000	239.268	250.909
<b>sqft_living15</b>	68.0150	3.962	17.168	0.000	60.250	75.780
<b>sqft_lot</b>	-0.3078	0.043	-7.085	0.000	-0.393	-0.223

<b>Omnibus:</b>	15375.814	<b>Durbin-Watson:</b>	1.989
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	666369.950
<b>Skew:</b>	2.967	<b>Prob(JB):</b>	0.00
<b>Kurtosis:</b>	29.672	<b>Cond. No.</b>	1.36e+05

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.36e+05. This might indicate that there are strong multicollinearity or other numerical problems.

In [66]:

```
Model_4 = ols(formula="price ~sqft_living + sqft_living15", data=df).fit()
Model_4.summary()
```

Out[66]:

## OLS Regression Results

<b>Dep. Variable:</b>	price	<b>R-squared:</b>	0.499
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.499
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	1.068e+04
<b>Date:</b>	Thu, 17 Feb 2022	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	15:32:09	<b>Log-Likelihood:</b>	-2.9749e+05
<b>No. Observations:</b>	21420	<b>AIC:</b>	5.950e+05
<b>Df Residuals:</b>	21417	<b>BIC:</b>	5.950e+05
<b>Df Model:</b>	2		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	-9.958e+04	5461.213	-18.234	0.000	-1.1e+05	-8.89e+04
<b>sqft_living</b>	243.0141	2.959	82.136	0.000	237.215	248.813
<b>sqft_living15</b>	67.4359	3.965	17.006	0.000	59.663	75.208

<b>Omnibus:</b>	15441.551	<b>Durbin-Watson:</b>	1.989
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	675536.703
<b>Skew:</b>	2.984	<b>Prob(JB):</b>	0.00
<b>Kurtosis:</b>	29.857	<b>Cond. No.</b>	9.43e+03

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 9.43e+03. This might indicate that there are strong multicollinearity or other numerical problems.

## POST MODEL ANALYSIS

In [ ]: